

Oil Production Forecasting Using Machine Learning

Mahammad Salman Shaik (2215886)

Akram Mohammad (2162967)

Imtiyaz Ali Syed (2220056)

Faisal Malik Mohammed (2214828)

Import necessary libraries and Load the data

```
In [1]: # Import necessary libraries
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import missingno as msno

# Load Volve production data from Excel file
volve_data = pd.read_excel('Volve production data.xlsx')
```

Performing Exploratory Data Analytics

```
In [2]: volve_data.head()
```

Out[2]:

	DATEPRD	WELL_BORE_CODE	NPD_WELL_BORE_CODE	NPD_WELL_BORE_NAME	NPD_FIE
0	2014-04-07	NO 15/9-F-1 C	7405	15/9-F-1 C	
1	2014-04-08	NO 15/9-F-1 C	7405	15/9-F-1 C	
2	2014-04-09	NO 15/9-F-1 C	7405	15/9-F-1 C	
3	2014-04-10	NO 15/9-F-1 C	7405	15/9-F-1 C	
4	2014-04-11	NO 15/9-F-1 C	7405	15/9-F-1 C	

5 rows × 24 columns

In [3]: *# Print information about the data, including column names, data types, and non-null values.*
`volve_data.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 15634 entries, 0 to 15633
Data columns (total 24 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   DATEPRD                              15634 non-null  datetime64[ns]
1   WELL_BORE_CODE                       15634 non-null  object
2   NPD_WELL_BORE_CODE                   15634 non-null  int64
3   NPD_WELL_BORE_NAME                   15634 non-null  object
4   NPD_FIELD_CODE                       15634 non-null  int64
5   NPD_FIELD_NAME                       15634 non-null  object
6   NPD_FACILITY_CODE                   15634 non-null  int64
7   NPD_FACILITY_NAME                   15634 non-null  object
8   ON_STREAM_HRS                        15349 non-null  float64
9   AVG_DOWNHOLE_PRESSURE                8980 non-null  float64
10  AVG_DOWNHOLE_TEMPERATURE              8980 non-null  float64
11  AVG_DP_TUBING                        8980 non-null  float64
12  AVG_ANNULUS_PRESS                    7890 non-null  float64
13  AVG_CHOKE_SIZE_P                     8919 non-null  float64
14  AVG_CHOKE_UOM                        9161 non-null  object
15  AVG_WHP_P                            9155 non-null  float64
16  AVG_WHT_P                            9146 non-null  float64
17  DP_CHOKE_SIZE                        15340 non-null  float64
18  BORE_OIL_VOL                         9161 non-null  float64
19  BORE_GAS_VOL                         9161 non-null  float64
20  BORE_WAT_VOL                         9161 non-null  float64
21  BORE_WI_VOL                          5706 non-null  float64
22  FLOW_KIND                            15634 non-null  object
23  WELL_TYPE                            15634 non-null  object
dtypes: datetime64[ns](1), float64(13), int64(3), object(7)
memory usage: 2.9+ MB
```

In [4]: *# We can check the shape and columns of the data*
`print(volve_data.shape)`
`print(volve_data.columns)`

```
(15634, 24)
Index(['DATEPRD', 'WELL_BORE_CODE', 'NPD_WELL_BORE_CODE', 'NPD_WELL_BORE_NAME',
      'NPD_FIELD_CODE', 'NPD_FIELD_NAME', 'NPD_FACILITY_CODE',
      'NPD_FACILITY_NAME', 'ON_STREAM_HRS', 'AVG_DOWNHOLE_PRESSURE',
      'AVG_DOWNHOLE_TEMPERATURE', 'AVG_DP_TUBING', 'AVG_ANNULUS_PRESS',
      'AVG_CHOKE_SIZE_P', 'AVG_CHOKE_UOM', 'AVG_WHP_P', 'AVG_WHT_P',
      'DP_CHOKE_SIZE', 'BORE_OIL_VOL', 'BORE_GAS_VOL', 'BORE_WAT_VOL',
      'BORE_WI_VOL', 'FLOW_KIND', 'WELL_TYPE'],
      dtype='object')
```

```
In [5]: # We can also check the data types of each column  
print(volve_data.dtypes)
```

```
DATEPRD                datetime64[ns]  
WELL_BORE_CODE          object  
NPD_WELL_BORE_CODE      int64  
NPD_WELL_BORE_NAME      object  
NPD_FIELD_CODE          int64  
NPD_FIELD_NAME          object  
NPD_FACILITY_CODE       int64  
NPD_FACILITY_NAME       object  
ON_STREAM_HRS           float64  
AVG_DOWNHOLE_PRESSURE   float64  
AVG_DOWNHOLE_TEMPERATURE float64  
AVG_DP_TUBING           float64  
AVG_ANNULUS_PRESS       float64  
AVG_CHOKE_SIZE_P        float64  
AVG_CHOKE_UOM           object  
AVG_WHP_P               float64  
AVG_WHT_P               float64  
DP_CHOKE_SIZE           float64  
BORE_OIL_VOL            float64  
BORE_GAS_VOL            float64  
BORE_WAT_VOL            float64  
BORE_WI_VOL             float64  
FLOW_KIND               object  
WELL_TYPE               object  
dtype: object
```

In [6]: `null_counts = volve_data.isnull().sum()`

```
# Print the result
print(null_counts)
```

```
DATEPRD                0
WELL_BORE_CODE          0
NPD_WELL_BORE_CODE      0
NPD_WELL_BORE_NAME      0
NPD_FIELD_CODE          0
NPD_FIELD_NAME          0
NPD_FACILITY_CODE       0
NPD_FACILITY_NAME       0
ON_STREAM_HRS           285
AVG_DOWNHOLE_PRESSURE   6654
AVG_DOWNHOLE_TEMPERATURE 6654
AVG_DP_TUBING           6654
AVG_ANNULUS_PRESS       7744
AVG_CHOKE_SIZE_P        6715
AVG_CHOKE_UOM           6473
AVG_WHP_P               6479
AVG_WHT_P               6488
DP_CHOKE_SIZE           294
BORE_OIL_VOL            6473
BORE_GAS_VOL            6473
BORE_WAT_VOL            6473
BORE_WI_VOL             9928
FLOW_KIND                0
WELL_TYPE                0
dtype: int64
```

In [7]: `# Generate summary statistics of the data, including count, mean, standard deviation`
`volve_data.describe()`

Out[7]:

	NPD_WELL_BORE_CODE	NPD_FIELD_CODE	NPD_FACILITY_CODE	ON_STREAM_HRS
count	15634.000000	15634.0	15634.0	15349.000000
mean	5908.581745	3420717.0	369304.0	19.994093
std	649.231622	0.0	0.0	8.369978
min	5351.000000	3420717.0	369304.0	0.000000
25%	5599.000000	3420717.0	369304.0	24.000000
50%	5693.000000	3420717.0	369304.0	24.000000
75%	5769.000000	3420717.0	369304.0	24.000000
max	7405.000000	3420717.0	369304.0	25.000000

```
In [8]: # Calculate the correlation between the different columns of the data
volve_data.corr()
```

Out[8]:


	NPD_WELL_BORE_CODE	NPD_FIELD_CODE	NPD_FACILITY_CODE
NPD_WELL_BORE_CODE	1.000000	NaN	NaN
NPD_FIELD_CODE	NaN	NaN	NaN
NPD_FACILITY_CODE	NaN	NaN	NaN
ON_STREAM_HRS	-0.102270	NaN	NaN
AVG_DOWNHOLE_PRESSURE	0.257481	NaN	NaN
AVG_DOWNHOLE_TEMPERATURE	0.339509	NaN	NaN
AVG_DP_TUBING	0.218243	NaN	NaN
AVG_ANNULUS_PRESS	0.141756	NaN	NaN
AVG_CHOKE_SIZE_P	-0.558461	NaN	NaN
AVG_WHP_P	0.077946	NaN	NaN
AVG_WHT_P	-0.519515	NaN	NaN
DP_CHOKE_SIZE	0.237647	NaN	NaN
BORE_OIL_VOL	-0.307645	NaN	NaN
BORE_GAS_VOL	-0.310793	NaN	NaN
BORE_WAT_VOL	-0.493591	NaN	NaN
BORE_WI_VOL	-0.055894	NaN	NaN

```
In [9]: # Count the number of occurrences of each value in the 'WELL_BORE_CODE' column
print(volve_data['WELL_BORE_CODE'].value_counts())
```

```
NO 15/9-F-4 AH      3327
NO 15/9-F-5 AH      3306
NO 15/9-F-12 H      3056
NO 15/9-F-14 H      3056
NO 15/9-F-11 H      1165
NO 15/9-F-15 D       978
NO 15/9-F-1 C       746
Name: WELL_BORE_CODE, dtype: int64
```

```
In [10]: # Use the missingno library to visualize the missing data in the dataset
missingno.matrix(volve_data)
```

Out[10]: <AxesSubplot:>



```
In [11]: print(volve_data['NPD_WELL_BORE_CODE'].value_counts())
```

```
5693    3327
5769    3306
5599    3056
5351    3056
7078    1165
7289     978
7405     746
Name: NPD_WELL_BORE_CODE, dtype: int64
```

```
In [12]: # Filter the data by selecting rows where the 'NPD_WELL_BORE_CODE' column matches
well_1_data = volve_data[volve_data['NPD_WELL_BORE_CODE'] == 7405]
well_2_data = volve_data[volve_data['NPD_WELL_BORE_CODE'] == 7078]
well_3_data = volve_data[volve_data['NPD_WELL_BORE_CODE'] == 5599]
well_4_data = volve_data[volve_data['NPD_WELL_BORE_CODE'] == 5351]
well_5_data = volve_data[volve_data['NPD_WELL_BORE_CODE'] == 7289]
well_6_data = volve_data[volve_data['NPD_WELL_BORE_CODE'] == 5693]
well_7_data = volve_data[volve_data['NPD_WELL_BORE_CODE'] == 5769]
```

```
In [13]: print("Information about well 1:")
print(well_1_data.info())
print("")

print("Information about well 2:")
print(well_2_data.info())
print("")

print("Information about well 3:")
print(well_3_data.info())
print("")

print("Information about well 4:")
print(well_4_data.info())
print("")

print("Information about well 5:")
print(well_5_data.info())
print("")

print("Information about well 6:")
print(well_6_data.info())
print("")

print("Information about well 7:")
print(well_7_data.info())
```

```
Information about well 1:
<class 'pandas.core.frame.DataFrame'>
Int64Index: 746 entries, 0 to 745
Data columns (total 24 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   DATEPRD                              746 non-null    datetime64[ns]
 1   WELL_BORE_CODE                       746 non-null    object
 2   NPD_WELL_BORE_CODE                   746 non-null    int64
 3   NPD_WELL_BORE_NAME                   746 non-null    object
 4   NPD_FIELD_CODE                       746 non-null    int64
 5   NPD_FIELD_NAME                       746 non-null    object
 6   NPD_FACILITY_CODE                    746 non-null    int64
 7   NPD_FACILITY_NAME                    746 non-null    object
 8   ON_STREAM_HRS                        746 non-null    float64
 9   AVG_DOWNHOLE_PRESSURE                743 non-null    float64
10  AVG_DOWNHOLE_TEMPERATURE              743 non-null    float64
11  AVG_DP_TUBING                        743 non-null    float64
12  AVG_ANNULUS_PRESS                    17 non-null     float64
13  AVG_CHUNK_SIZE                       746 non-null    int64
```



```

In [14]: import numpy as np
from matplotlib import pyplot as plt
%matplotlib inline
plt.style.use('ggplot')

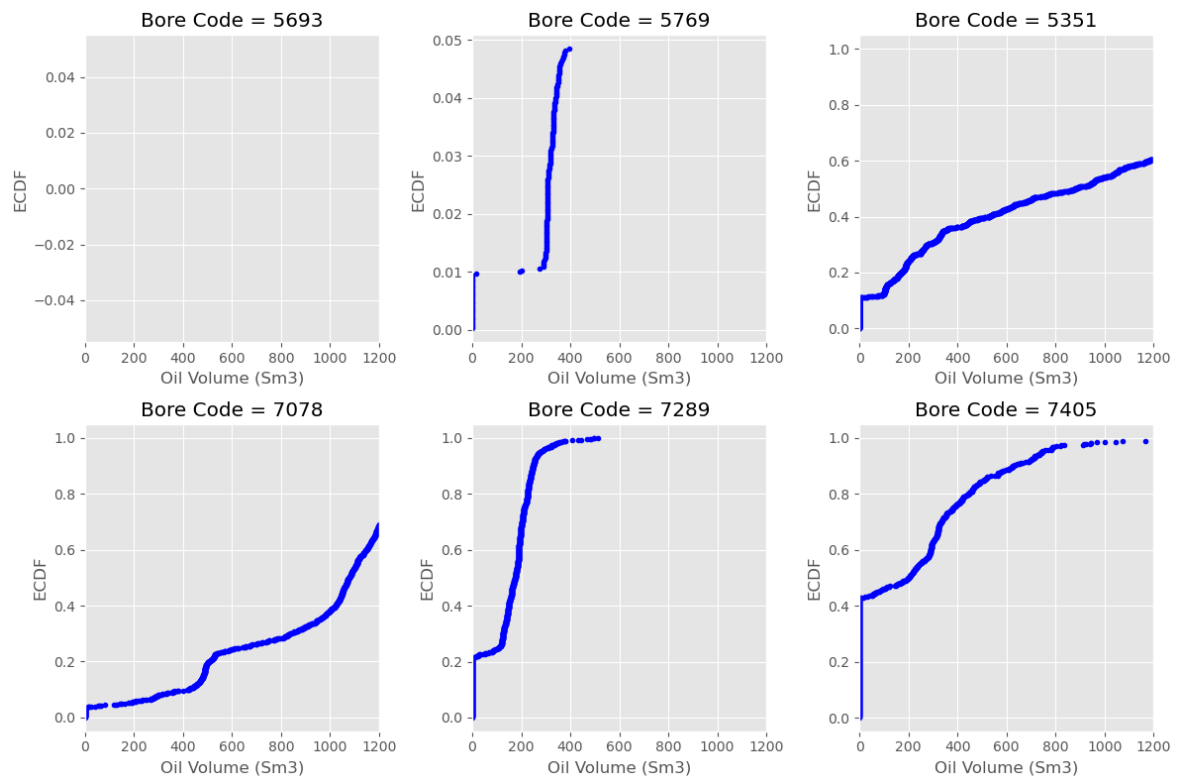
# Define function for creating ECDF plot
def ecdf(data):
    n = len(data)
    x = np.sort(data)
    y = np.arange(1,n+1)/n
    return x,y

# Create subplots for all wells
fig, axs = plt.subplots(2, 3, figsize=(12, 8))
axs = axs.ravel() # flatten the array of subplots

# Loop through wells and plot ECDF
wells = [5693, 5769, 5351, 7078, 7289, 7405]
for i, well in enumerate(wells):
    well_data = volve_data[volve_data['NPD_WELL_BORE_CODE'] == well]
    x_axis, y_axis = ecdf(well_data['BORE_OIL_VOL'])
    axs[i].plot(x_axis, y_axis, marker=".", linestyle="none", color='blue')
    axs[i].set_xlabel('Oil Volume (Sm3)')
    axs[i].set_ylabel('ECDF')
    axs[i].set_title(f'Bore Code = {well}')
    axs[i].set_xlim([0, 1200]) # set limit for x-axis

plt.tight_layout()
plt.show()


```



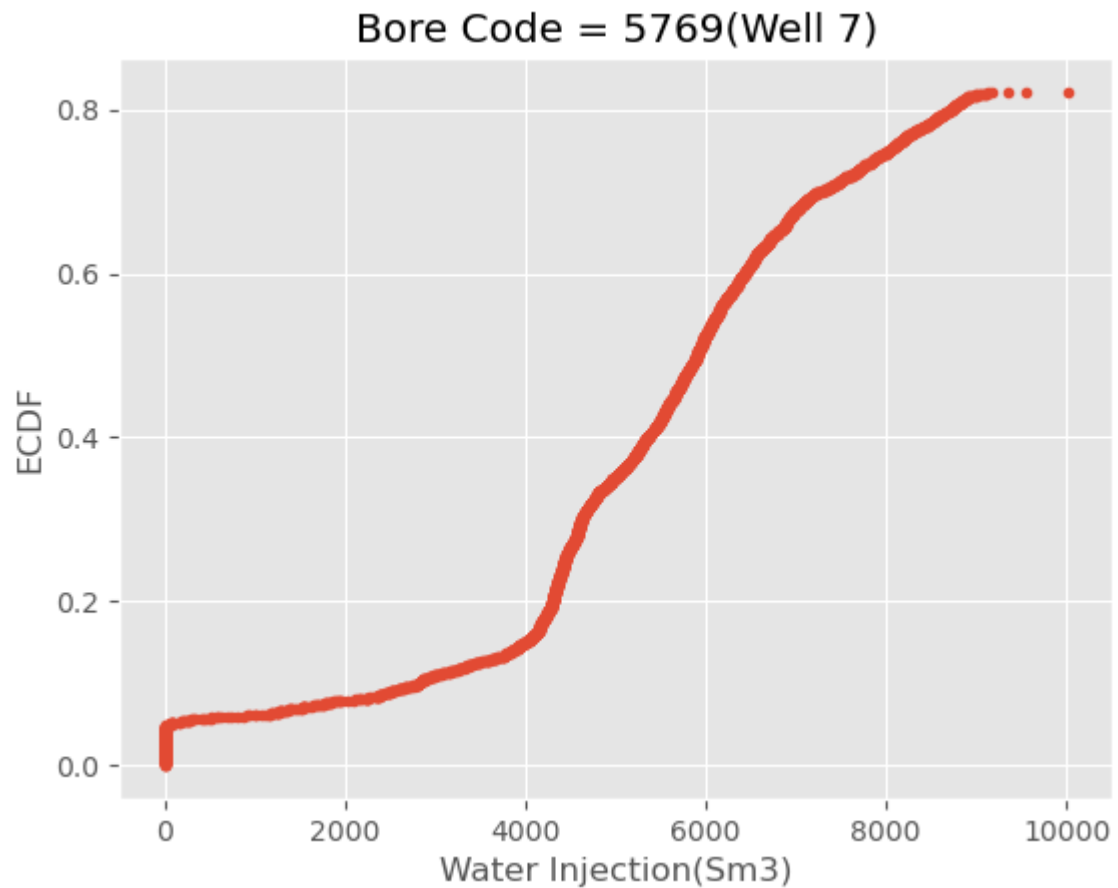
```
In [15]: x_axis, y_axis = ecdf(well_6_data['BORE_WI_VOL'])
plt.plot(x_axis, y_axis, marker=".", linestyle="none")

#Labeling
plt.xlabel('Water Injection(Sm3)')
plt.ylabel('ECDF')
plt.title('Bore Code = 5693(Well 6)')

plt.show()
```



```
In [16]: x_axis, y_axis = ecdf(well_7_data['BORE_WI_VOL'])  
plt.plot(x_axis, y_axis, marker=".", linestyle="none")  
  
#Labeling  
plt.xlabel('Water Injection(Sm3)')  
plt.ylabel('ECDF')  
plt.title('Bore Code = 5769(Well 7)')  
  
plt.show()
```



```

In [17]: import seaborn as sns
import matplotlib.pyplot as plt

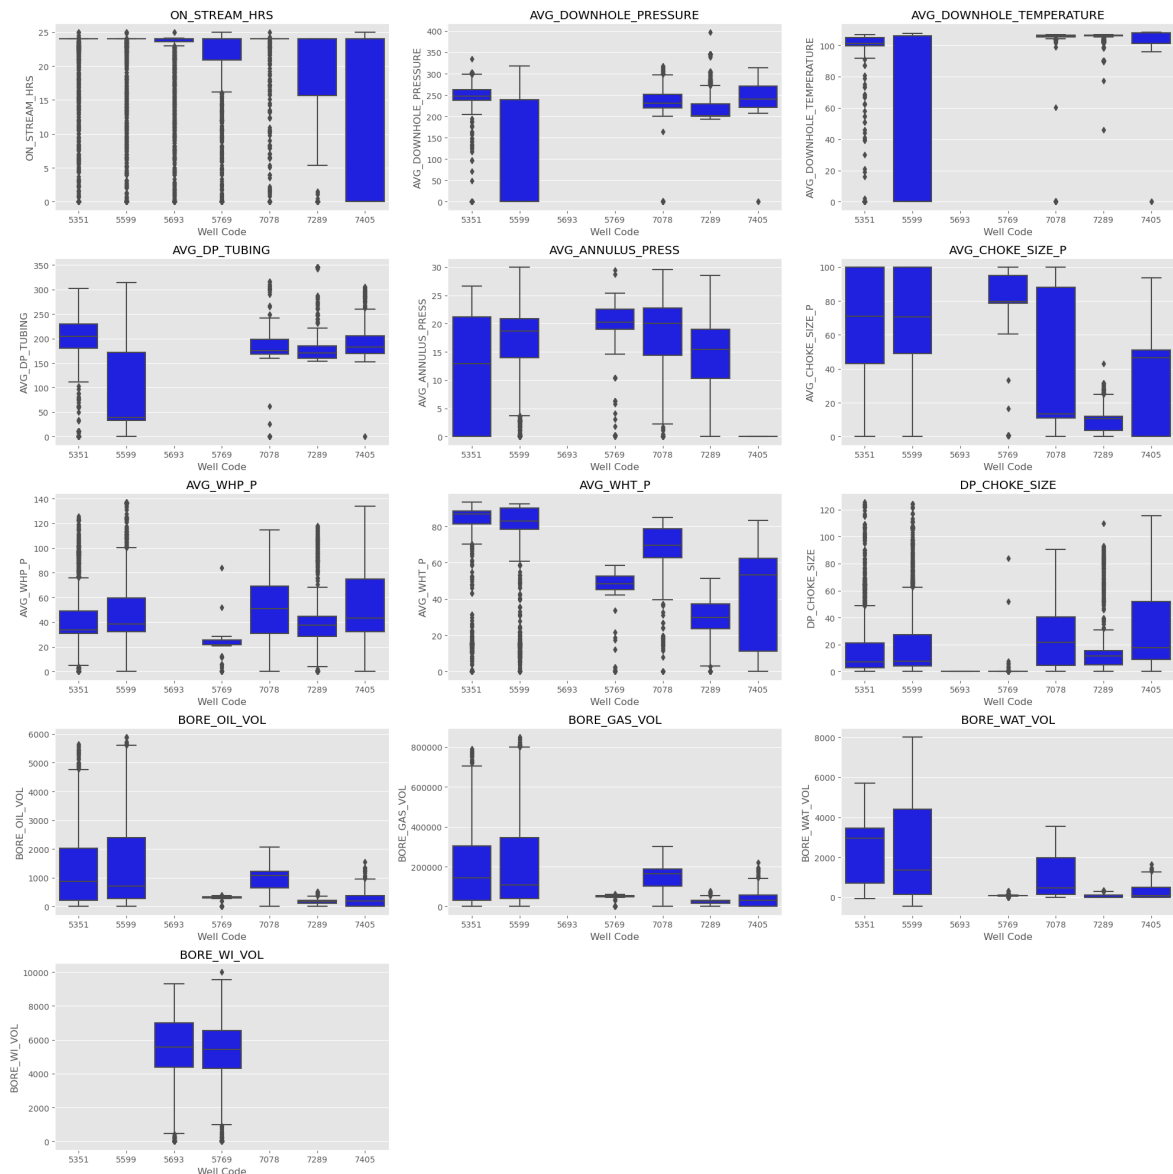
plt.figure(figsize=(20,20))

columns = ['ON_STREAM_HRS', 'AVG_DOWNHOLE_PRESSURE', 'AVG_DOWNHOLE_TEMPERATURE',
           'AVG_DP_TUBING', 'AVG_ANNULUS_PRESS', 'AVG_CHOKE_SIZE_P',
           'AVG_WHP_P', 'AVG_WHT_P', 'DP_CHOKE_SIZE', 'BORE_OIL_VOL',
           'BORE_GAS_VOL', 'BORE_WAT_VOL', 'BORE_WI_VOL']

for i, col in enumerate(columns):
    plt.subplot(5, 3, i+1)
    sns.boxplot(x='NPD_WELL_BORE_CODE', y=col, data=volve_data, color='blue')
    plt.xlabel('Well Code')
    plt.ylabel(col)
    plt.title(col)

plt.tight_layout()
plt.show()

```



```

In [18]: import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

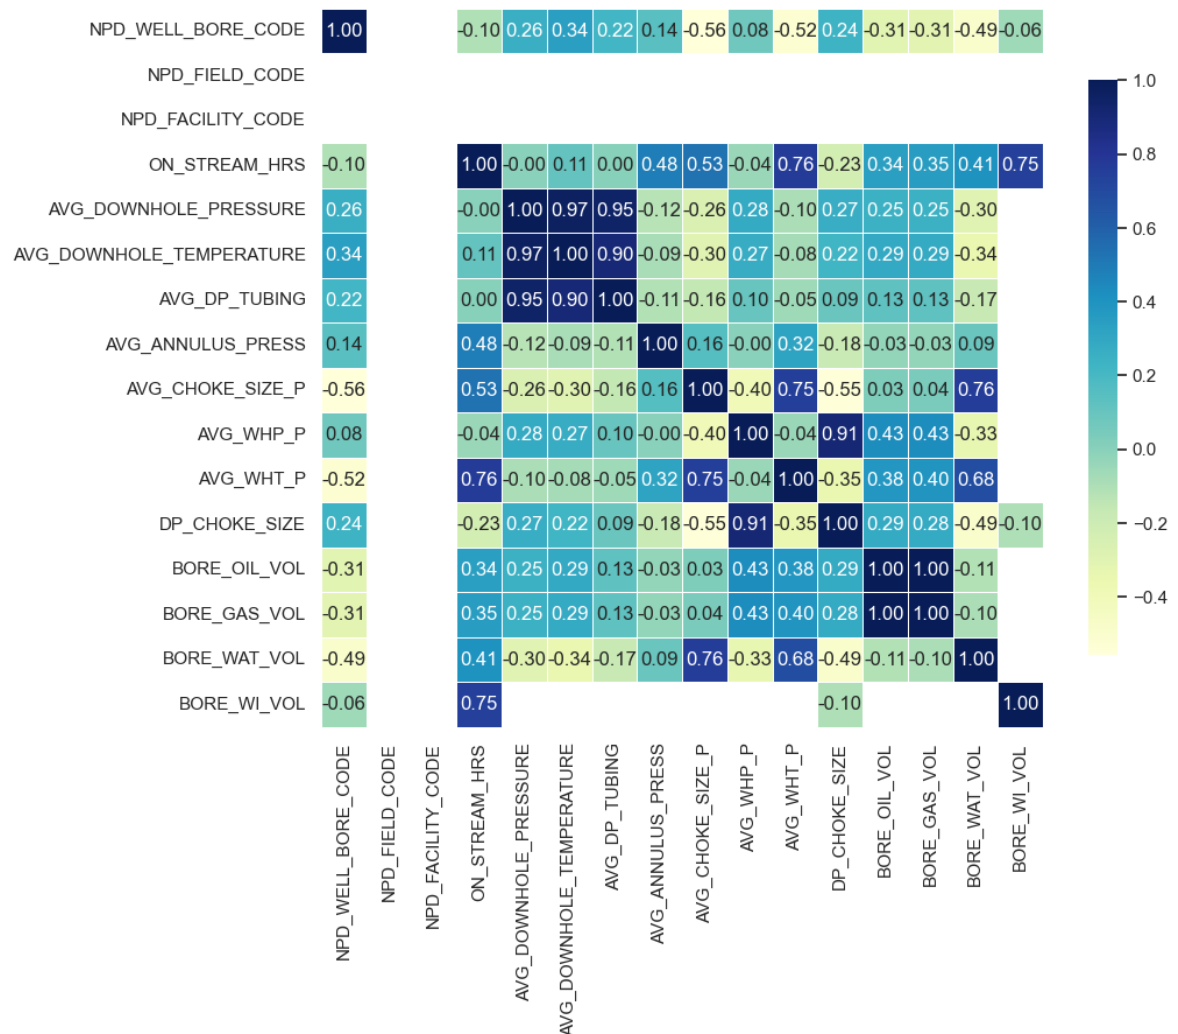
# Calculate the correlation matrix
corr_matrix = volve_data.corr()

# Set up figure size and style
sns.set(style="white")
plt.figure(figsize=(10, 8))

# Create correlation plot
sns.heatmap(corr_matrix, cmap='YlGnBu', annot=True, fmt='.2f', linewidths=.5,

```

Out[18]: <AxesSubplot:>



Performing Linear Regression

Import necessary libraries

```
In [19]: import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LinearRegression
from sklearn.preprocessing import MinMaxScaler
from sklearn.metrics import r2_score, mean_absolute_error, mean_squared_error
import math
from matplotlib import pyplot as plt
```

Load data

```
In [20]: df = pd.read_excel('Volve production data.xlsx')
```

Round data to nearest integer

```
In [21]: df = np.round(df)
```

Interpolate missing values

```
In [22]: cols_to_interpolate = ['ON_STREAM_HRS', 'AVG_DOWNHOLE_PRESSURE', 'AVG_DOWNHOL
    'AVG_ANNULUS_PRESS', 'AVG_CHOKE_SIZE_P', 'AVG_WHP_P',
    'BORE_OIL_VOL', 'BORE_GAS_VOL', 'BORE_WAT_VOL']
df[cols_to_interpolate] = df[cols_to_interpolate].interpolate(method='linear')
```

Drop unnecessary columns

```
In [23]: cols_to_drop = ['WELL_BORE_CODE', 'NPD_WELL_BORE_NAME', 'NPD_FIELD_NAME', 'NPI
    'NPD_FACILITY_NAME', 'AVG_DOWNHOLE_PRESSURE', 'AVG_DP_TUBING'
    'FLOW_KIND', 'WELL_TYPE', 'BORE_GAS_VOL']
df.drop(cols_to_drop, axis=1, inplace=True)
```

Scale dataset

```
In [24]: cols_to_scale = ['ON_STREAM_HRS', 'AVG_DOWNHOLE_TEMPERATURE', 'AVG_ANNULUS_PR',  
                        'AVG_WHT_P']  
scaler = MinMaxScaler()  
df[cols_to_scale] = scaler.fit_transform(df[cols_to_scale])
```

Select data for wells 5599, 5351, 7078

```
In [25]: df = df.loc[df['NPD_WELL_BORE_CODE'].isin([5599, 5351, 7078])]
```

Split data into training and testing sets

```
In [26]: X = df.drop(['BORE_OIL_VOL'],axis=1)  
y= df[['BORE_OIL_VOL','NPD_WELL_BORE_CODE']]  
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Train a linear regression model

```
In [27]: reg_linear = LinearRegression()  
x_train_final = X_train.drop(['DATEPRD','NPD_WELL_BORE_CODE'],axis = 1)  
x_test_final = X_test.drop(['DATEPRD','NPD_WELL_BORE_CODE'],axis = 1)  
y_test_final = y_test['BORE_OIL_VOL']  
y_train_final = y_train['BORE_OIL_VOL']  
reg_linear.fit(x_train_final, y_train_final)
```

```
Out[27]: LinearRegression()
```

Use the model to predict oil production for the test data

```
In [28]: y_pred_linear = reg_linear.predict(x_test_final)
```

Calculate the Performance Metrics

```
In [29]: # Calculate R-squared score to evaluate the model's performance
r2 = r2_score(y_test_final, y_pred_linear)
print('R-squared score of Linear Regression Model is :', r2)

# Calculate mean absolute error
mae = mean_absolute_error(y_test_final, y_pred_linear)
print('Mean absolute error of Linear Regression Model is :', mae)


# Calculate root mean squared error
rmse = math.sqrt(mean_squared_error(y_test_final, y_pred_linear))
print('Root mean squared error of linear Regression Model:', rmse)
# Calculate mean squared error
mse = mean_squared_error(y_test_final, y_pred_linear)
print('mean squared error of Linear Regression Model is :', mse)
```

R-squared score of Linear Regression Model is : 0.554546213401562
Mean absolute error of Linear Regression Model is : 705.8758754656305
Root mean squared error of linear Regression Model: 913.4242223150642
mean squared error of Linear Regression Model is : 834343.8099118797

Plot actual vs predicted values for Linear Regression Model

```
In [30]: # Plot actual vs predicted values
plt.style.use('fivethirtyeight')
plt.figure(figsize = (14,8))
plt.scatter(X_test["DATEPRD"].tolist(), y_pred_linear, label='predicted')
plt.scatter(X_test["DATEPRD"].tolist(), y_test_final, label='actual')
plt.legend()
plt.xlabel("Year")
plt.ylabel("Bore oil volume")
plt.title('Actual v/s Linear Regression Model prediction for Bore Oil Volume')

# Show plot
plt.show()
```



Plotting the actual vs Predicted values for Each Well Individually



Plot actual vs predicted values for 5599

```

In [31]: # Select data for well 5599
X_test_5599 = X_test[X_test["NPD_WELL_BORE_CODE"] == 5599]
y_test_5599 = y_test[y_test['NPD_WELL_BORE_CODE'] == 5599]
x_test_5599final = X_test_5599.drop(['DATEPRD', "NPD_WELL_BORE_CODE"], axis = 1)


# Predict values
y_pred_linear = reg_linear.predict(x_test_5599final)

# Plot actual vs predicted values
plt.style.use('fivethirtyeight')
plt.figure(figsize = (14,8))
plt.scatter(X_test_5599["DATEPRD"].tolist(), y_pred_linear, label='predicted')
plt.scatter(X_test_5599["DATEPRD"].tolist(), y_test_5599['BORE_OIL_VOL'], label='actual')
plt.legend()
plt.xlabel("Year")
plt.ylabel("Bore oil volume")
plt.title('Actual v/s Linear Regression Model prediction for Bore Oil Volume for Well No. - 5599')

# Evaluate performance
y_test_5599_final = y_test_5599['BORE_OIL_VOL']
print("The R2 value for linear regression for oil volume production in well 5599 is")

```

The R2 value for linear regression for oil volume production in well 5599 is 0.550247303910232



Plotting Actual vs predicted values for Well 5351

```

In [32]: # Select data for well 5351
X_test_5351 = X_test[X_test["NPD_WELL_BORE_CODE"] == 5351]
y_test_5351 = y_test[y_test['NPD_WELL_BORE_CODE'] == 5351]
x_test_5351final = X_test_5351.drop(['DATEPRD', "NPD_WELL_BORE_CODE"], axis = 1)


# Predict values
y_pred_linear = reg_linear.predict(x_test_5351final)

# Plot actual vs predicted values
plt.style.use('fivethirtyeight')
plt.figure(figsize = (14,8))
plt.scatter(X_test_5351["DATEPRD"].tolist(), y_pred_linear, label='predicted')
plt.scatter(X_test_5351["DATEPRD"].tolist(), y_test_5351['BORE_OIL_VOL'], label='actual')
plt.legend()
plt.xlabel("Year")
plt.ylabel("Bore oil volume")
plt.title('Actual v/s Linear Regression Model prediction for Bore Oil Volume for Well No. - 5351')

# Evaluate performance
y_test_5351_final = y_test_5351['BORE_OIL_VOL']
print("The R2 value for linear regression for oil volume production in well 5351 is")

```

The R2 value for linear regression for oil volume production in well 5351 is 0.6742749337876078



Plotting Actual vs predicted values for Well 7078

```


In [33]: # Select data for well 7078
X_test_7078 = X_test[X_test["NPD_WELL_BORE_CODE"] == 7078]
y_test_7078 = y_test[y_test['NPD_WELL_BORE_CODE'] == 7078]
x_test_7078final = X_test_7078.drop(['DATEPRD', "NPD_WELL_BORE_CODE"], axis = 1)

# Predict using the trained model
y_pred_linear = reg_linear.predict(x_test_7078final)

# Plot actual vs predicted values
plt.style.use('fivethirtyeight')
plt.figure(figsize = (14,8))
plt.scatter(X_test_7078["DATEPRD"].tolist(), y_pred_linear, label='predicted')
plt.scatter(X_test_7078["DATEPRD"].tolist(), y_test_7078['BORE_OIL_VOL'], label='actual')
plt.legend()
plt.xlabel("Year")
plt.ylabel("Bore oil volume")
plt.title('Actual v/s Linear Regression Model prediction for Bore Oil Volume for Well No. - 7078')

```

Out[33]: Text(0.5, 1.0, 'Actual v/s Linear Regression Model prediction for Bore Oil Volume for Well No. - 7078')



Polynomial regression

Importing any necessary libraries

```

In [34]: from sklearn.preprocessing import PolynomialFeatures
from sklearn.linear_model import LinearRegression

```

Create polynomial features of degree 2

```
In [35]: poly = PolynomialFeatures(degree=2)
x_train_poly = poly.fit_transform(x_train_final)
x_test_poly = poly.transform(x_test_final)
```

Fit linear regression model

```
In [36]: reg_poly = LinearRegression()
reg_poly.fit(x_train_poly, y_train_final)
```

```
Out[36]: LinearRegression()
```

Predict on test data

```
In [37]: y_pred_poly = reg_poly.predict(x_test_poly)
```

Calculating Performance metrics for Polynomial Regression

```
In [38]: # Calculate R-squared score
r2_poly = r2_score(y_test_final, y_pred_poly)
print('R-squared score of Ploynomial Regression is :', r2_poly)

# Calculate mean absolute error
mae_poly = mean_absolute_error(y_test_final, y_pred_poly)
print('Mean absolute error of Ploynomial Regression is :', mae_poly)


#Calculate mean squared error
mse_poly = mean_squared_error(y_test_final, y_pred_poly)
print('Mean squared error of Ploynomial Regression is :', mse_poly)
# Calculate root mean squared error
rmse_poly = math.sqrt(mean_squared_error(y_test_final, y_pred_poly))
print('Root mean squared error of Ploynomial Regression is :', rmse_poly)
```

```
R-squared score of Ploynomial Regression is : 0.9301328617036829
Mean absolute error of Ploynomial Regression is : 250.37863563033665
Mean squared error of Ploynomial Regression is : 130862.54086855218
Root mean squared error of Ploynomial Regression is : 361.7492790159397
```

Plotting Actual vs Predicted values for Polynomial Regression

```
In [39]: # Plot actual vs predicted values
plt.style.use('fivethirtyeight')
#plt.style.use('seaborn')
plt.figure(figsize=(14, 8))
plt.scatter(X_test['DATEPRD'], y_pred_poly, label='predicted')
plt.scatter(X_test['DATEPRD'], y_test_final, label='actual')
plt.legend()
plt.xlabel("Year")
plt.ylabel("Bore oil volume")
plt.title('Actual v/s Polynomial Regression Model prediction for Bore Oil Volume')

# Show plot
plt.show()
```




Plotting Actual vs predicted values for each Well

Plotting for well 5599

```
In [40]: # Perform polynomial regression for well 5599
poly = PolynomialFeatures(degree=2)
x_poly_5599 = poly.fit_transform(x_test_5599final)
reg_poly = LinearRegression()
reg_poly.fit(x_poly_5599, y_test_5599_final)
y_pred_poly = reg_poly.predict(poly.fit_transform(x_test_5599final))

# Plot actual vs predicted values
plt.figure(figsize=(14, 8))
plt.scatter(X_test_5599["DATEPRD"].tolist(), y_pred_poly, label='Polynomial Regression')
plt.scatter(X_test_5599["DATEPRD"].tolist(), y_test_5599['BORE_OIL_VOL'], label='Actual')
plt.legend()
plt.xlabel("Year")
plt.ylabel("Bore oil volume")
plt.title('Actual v/s Polynomial Regression Model prediction for Bore Oil Volume for Well No. - 5599')
```

Out[40]: Text(0.5, 1.0, 'Actual v/s Polynomial Regression Model prediction for Bore Oil Volume for Well No. - 5599')



Plotting for Well 5351


```

In [41]: # Perform polynomial regression for well 5351
poly = PolynomialFeatures(degree=2)
x_poly_5351 = poly.fit_transform(x_test_5351final)
reg_poly = LinearRegression()
reg_poly.fit(x_poly_5351, y_test_5351_final)
y_pred_poly = reg_poly.predict(poly.fit_transform(x_test_5351final))

# Plot actual vs predicted values
plt.figure(figsize=(14, 8))
plt.scatter(X_test_5351["DATEPRD"].tolist(), y_pred_poly, label='Polynomial Regression')
plt.scatter(X_test_5351["DATEPRD"].tolist(), y_test_5351['BORE_OIL_VOL'], label='Actual')
plt.legend()
plt.xlabel("Year")
plt.ylabel("Bore oil volume")
plt.title('Actual v/s Polynomial Regression Model prediction for Bore Oil Volume for Well No. - 5351')

```

Out[41]: Text(0.5, 1.0, 'Actual v/s Polynomial Regression Model prediction for Bore Oil Volume for Well No. - 5351')




Plotting for Well 7078


```
In [42]: poly = PolynomialFeatures(degree=2)
y_test_7078_final = y_test_7078['BORE_OIL_VOL']
x_poly_7078 = poly.fit_transform(x_test_7078final)
reg_poly = LinearRegression()
reg_poly.fit(x_poly_7078, y_test_7078_final)
y_pred_poly = reg_poly.predict(poly.fit_transform(x_test_7078final))

# Plot actual vs predicted values
plt.figure(figsize=(14, 8))
plt.scatter(X_test_7078["DATEPRD"].tolist(), y_pred_poly, label='Polynomial Regression')
plt.scatter(X_test_7078["DATEPRD"].tolist(), y_test_7078['BORE_OIL_VOL'], label='Actual')
plt.legend()
plt.xlabel("Year")
plt.ylabel("Bore oil volume")
plt.title('Actual v/s Polynomial Regression Model prediction for Bore Oil Volume for Well No. - 7078')
```

Out[42]: Text(0.5, 1.0, 'Actual v/s Polynomial Regression Model prediction for Bore Oil Volume for Well No. - 7078')



Type *Markdown* and LaTeX: α^2

XGBOOST

Import the Libraries

```
In [43]: import pandas as pd
import numpy as np
import xgboost as xgb
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

Load the data

```
In [44]: df = pd.read_excel('Volve production data.xlsx')
```

Round data to nearest integer

```
In [45]: # Round data to nearest integer
df = np.round(df)
```

Convert date to Unix timestamp

```
In [46]: # Convert date to Unix timestamp
df['DATEPRD'] = pd.to_datetime(df['DATEPRD'])
df['DATEPRD'] = (df['DATEPRD'] - pd.Timestamp("1970-01-01")) // pd.Timedelta(
```

Interpolate missing values

```
In [47]: # Interpolate missing values
cols_to_interpolate = ['ON_STREAM_HRS', 'AVG_DOWNHOLE_PRESSURE', 'AVG_DOWNHOL
                    'AVG_ANNULUS_PRESS', 'AVG_CHOKE_SIZE_P', 'AVG_WHP_P',
                    'BORE_OIL_VOL', 'BORE_GAS_VOL', 'BORE_WAT_VOL']
df[cols_to_interpolate] = df[cols_to_interpolate].interpolate(method='linear')
```

Drop unnecessary columns

```
In [48]: # Drop unnecessary columns
cols_to_drop = ['WELL_BORE_CODE', 'NPD_WELL_BORE_NAME', 'NPD_FIELD_NAME', 'NPD_FACILITY_NAME', 'AVG_DOWNHOLE_PRESSURE', 'AVG_DP_TUBING', 'FLOW_KIND', 'WELL_TYPE', 'BORE_GAS_VOL']
df.drop(cols_to_drop, axis=1, inplace=True)
```

Select data for wells 5599,5351,7078

```
In [49]: # Select data for wells 2-4
df = df.loc[df['NPD_WELL_BORE_CODE'].isin([5599,5351,7078])]
```

Prepare data for training

```
In [50]: # Prepare data for training
X = df.drop(['BORE_OIL_VOL'], axis=1)
y = df['BORE_OIL_VOL']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_state=42)
```

Train model

```
In [51]: # Train model
model = xgb.XGBRegressor(n_estimators=100, learning_rate=0.1, gamma=0, subsample=0.5,
                          colsample_bytree=1, max_depth=7, random_state=42)
model.fit(X_train, y_train)
```

```
Out[51]: XGBRegressor(base_score=None, booster=None, callbacks=None,
                      colsample_bylevel=None, colsample_bynode=None, colsample_bytree=1,
                      early_stopping_rounds=None, enable_categorical=False,
                      eval_metric=None, feature_types=None, gamma=0, gpu_id=None,
                      grow_policy=None, importance_type=None,
                      interaction_constraints=None, learning_rate=0.1, max_bin=None,
                      max_cat_threshold=None, max_cat_to_onehot=None,
                      max_delta_step=None, max_depth=7, max_leaves=None,
                      min_child_weight=None, missing=nan, monotone_constraints=None,
                      n_estimators=100, n_jobs=None, num_parallel_tree=None,
                      predictor=None, random_state=42, ...)
```

Predict on test set and calculate performance metrics

```
In [52]: # Predict on test set and calculate performance metrics
y_pred_xgb = model.predict(X_test)
r2 = r2_score(y_test, y_pred_xgb)
mae = mean_absolute_error(y_test, y_pred_xgb)
mse = mean_squared_error(y_test, y_pred_xgb)
rmse = np.sqrt(mse)
mape = np.mean(np.abs((y_test.replace(0, 1e-8) - y_pred_xgb) / y_test.replace(0, 1e-8)))

print("The R2 value for XGBoost Regression for oil volume production is", r2)
print("The mean absolute error (MAE) for XGBoost Regression is", mae)
print("The mean squared error (MSE) for XGBoost Regression is", mse)
print("The root mean squared error (RMSE) for XGBoost Regression is", rmse)
```

The R2 value for XGBoost Regression for oil volume production is 0.9940663288525774


The mean absolute error (MAE) for XGBoost Regression is 53.5111157297474

The mean squared error (MSE) for XGBoost Regression is 11113.884180240848

The root mean squared error (RMSE) for XGBoost Regression is 105.4224083401667

```
In [53]: # Plot actual vs predicted values
plt.style.use('fivethirtyeight')
#plt.style.use('ggplot')
plt.figure(figsize=(14, 8))
plt.scatter(X_test['DATEPRD'], y_pred_xgb, label='predicted')
plt.scatter(X_test['DATEPRD'], y_test_final, label='actual')
plt.legend()
plt.xlabel("Year ")#represented in terms of 9th power of seconds as we are us
plt.ylabel("Bore oil volume")
plt.title('Actual v/s XG Boost Regression Model prediction for Bore Oil Volume')

# Show plot
plt.show()
```



Plotting the Actual vs Predicted values for the wells using XGBoost Regression Model


Actual vs Predicted values for the well-5599


```
In [54]: import matplotlib.pyplot as plt

# Select data for well 7405
X_test_5599 = X_test[X_test['NPD_WELL_BORE_CODE'] == 5599]
y_test_5599 = y_test[X_test['NPD_WELL_BORE_CODE'] == 5599]
y_pred_5599 = model.predict(X_test_5599)

# Plot predicted vs actual oil production
plt.figure(figsize=(14, 8))
plt.scatter(X_test_5599['DATEPRD'], y_test_5599, label='actual')
plt.scatter(X_test_5599['DATEPRD'], y_pred_5599, label='predicted')
plt.xlabel('Year')
plt.ylabel('Bore oil volume')
plt.title('Actual vs XGBoost Regression Model prediction for Bore Oil Volume')
plt.legend()
plt.show()

# Plot residual plot
plt.figure(figsize=(14, 8))
plt.scatter(y_pred_5599, y_pred_5599 - y_test_5599, c='b', s=40, alpha=0.5)
plt.hlines(y=0, xmin=0, xmax=max(y_test_5599), colors='r', zorder=3)
plt.xlabel('Predicted values')
plt.ylabel('Residuals')
plt.title('Residual plot for XGBoost Regression Model for Well No. - 5599')
plt.show()
```





Actual vs Predicted values for the well-5351


```


In [55]: # Select data for well 5351
X_test_5351 = X_test[X_test['NPD_WELL_BORE_CODE'] == 5351]
y_test_5351 = y_test[X_test['NPD_WELL_BORE_CODE'] == 5351]
y_pred_5351 = model.predict(X_test_5351)

# Plot predicted vs actual oil production
plt.figure(figsize=(14, 8))
plt.scatter(X_test_5351['DATEPRD'], y_test_5351, label='actual')
plt.scatter(X_test_5351['DATEPRD'], y_pred_5351, label='predicted')
plt.xlabel('Year')
plt.ylabel('Bore oil volume')
plt.title('Actual vs XGBoost Regression Model prediction for Bore Oil Volume')
plt.legend()
plt.show()

# Plot residual plot
plt.figure(figsize=(14, 8))
plt.scatter(y_pred_5351, y_pred_5351 - y_test_5351, c='b', s=40, alpha=0.5)
plt.hlines(y=0, xmin=0, xmax=max(y_test_5351), colors='r', zorder=3)
plt.xlabel('Predicted values')
plt.ylabel('Residuals')
plt.title('Residual plot for XGBoost Regression Model for Well No. - 5351')
plt.show()

```





Actual vs Predicted values for the well-7078


```

In [56]: # Select data for well 7078
X_test_7078 = X_test[X_test['NPD_WELL_BORE_CODE'] == 7078]
y_test_7078 = y_test[X_test['NPD_WELL_BORE_CODE'] == 7078]
y_pred_7078 = model.predict(X_test_7078)

# Plot predicted vs actual oil production
plt.figure(figsize=(14, 8))
plt.scatter(X_test_7078['DATEPRD'], y_test_7078, label='actual')
plt.scatter(X_test_7078['DATEPRD'], y_pred_7078, label='predicted')
plt.xlabel('Year')
plt.ylabel('Bore oil volume')
plt.title('Actual vs XGBoost Regression Model prediction for Bore Oil Volume')
plt.legend()
plt.show()

# Plot residual plot
plt.figure(figsize=(14, 8))
plt.scatter(y_pred_7078, y_pred_7078 - y_test_7078, c='b', s=40, alpha=0.5)
plt.hlines(y=0, xmin=0, xmax=max(y_test_7078), colors='r', zorder=3)


```



```

Out[56]: <matplotlib.collections.LineCollection at 0x1f3846f3a30>

```



Random Forest

Importing the necessary libraries

```
In [57]: import pandas as pd
import numpy as np
from sklearn.ensemble import RandomForestRegressor
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_absolute_error, mean_squared_error, r2_score
```

Load data

```
In [58]: # Load data
df = pd.read_excel('Volve production data.xlsx')
```

Round data to nearest integer

```
In [59]: df = np.round(df)
```

Convert date to Unix timestamp

```
In [60]: df['DATEPRD'] = pd.to_datetime(df['DATEPRD'])
df['DATEPRD'] = (df['DATEPRD'] - pd.Timestamp("1970-01-01")) // pd.Timedelta(
```

Interpolate missing values

```
In [61]: cols_to_interpolate = ['ON_STREAM_HRS', 'AVG_DOWNHOLE_PRESSURE', 'AVG_DOWNHOL
      'AVG_ANNULUS_PRESS', 'AVG_CHOKE_SIZE_P', 'AVG_WHP_P',
      'BORE_OIL_VOL', 'BORE_GAS_VOL', 'BORE_WAT_VOL']
df[cols_to_interpolate] = df[cols_to_interpolate].interpolate(method='linear')
```

Drop unnecessary columns

```
In [62]: cols_to_drop = ['WELL_BORE_CODE', 'NPD_WELL_BORE_NAME', 'NPD_FIELD_NAME', 'NPD
      'NPD_FACILITY_NAME', 'AVG_DOWNHOLE_PRESSURE', 'AVG_DP_TUBING'
      'FLOW_KIND', 'WELL_TYPE', 'BORE_GAS_VOL']
df.drop(cols_to_drop, axis=1, inplace=True)
```

Select data for wells 5351,5599,7078

```
In [63]: df = df.loc[df['NPD_WELL_BORE_CODE'].isin([5351, 7078, 5599])]
```

Prepare data for training

```
In [64]: # Prepare data for training
X = df.drop(['BORE_OIL_VOL'], axis=1)
y = df['BORE_OIL_VOL']

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.4, random
```

Train model

```
In [65]: rf_model = RandomForestRegressor(n_estimators=100, max_depth=7, random_state=42)
rf_model.fit(X_train, y_train)
```

```
Out[65]: RandomForestRegressor(max_depth=7, random_state=42)
```

Predict on test set and calculate performance metrics

```
In [66]: y_pred = rf_model.predict(X_test)
r2 = r2_score(y_test, y_pred)
mae = mean_absolute_error(y_test, y_pred)
mse = mean_squared_error(y_test, y_pred)
rmse = np.sqrt(mse)
mape = np.mean(np.abs((y_test.replace(0, 1e-8) - y_pred) / y_test.replace(0,

print("The R2 value for Random Forest Regression for oil volume production is
print("The mean absolute error (MAE) for Random Forest Regression is", mae)
print("The mean squared error (MSE) for Random Forest Regression is", mse)
print("The root mean squared error (RMSE) for Random Forest Regression is", r
```

The R2 value for Random Forest Regression for oil volume production is 0.9755265282036907

The mean absolute error (MAE) for Random Forest Regression is 122.83331707517738

The mean squared error (MSE) for Random Forest Regression is 45689.43330414476


The root mean squared error (RMSE) for Random Forest Regression is 213.75086737635652

Plotting the actual vs Predicted values for random Forest Model

```
In [67]: import matplotlib.pyplot as plt

# Plot actual vs predicted values
plt.style.use('fivethirtyeight')
#plt.style.use('dark_background')
plt.figure(figsize=(14, 8))
plt.scatter(X_test['DATEPRD'], y_pred, label='predicted')
plt.scatter(X_test['DATEPRD'], y_test, label='actual')
plt.legend()
plt.xlabel("Year ")#represented in terms of 9th power of seconds as we are us
plt.ylabel("Bore oil volume")
plt.title('Actual vs Predicted Oil Volume Production (Random Forest Regression)')

# Show plot
plt.show()
```



In []:

In []:

In []:

Plotting the actual vs Predicted values for Each wells

Actual VS Predicted values for well-5351


```


In [68]: # Select data for well 5351
X_test_5351_rf = X_test[X_test['NPD_WELL_BORE_CODE'] == 5351]
y_test_5351_rf = y_test[X_test['NPD_WELL_BORE_CODE'] == 5351]
y_pred_5351_rf = rf_model.predict(X_test_5351_rf)

# Plot predicted vs actual oil production
plt.figure(figsize=(14, 8))
plt.scatter(X_test_5351_rf['DATEPRD'], y_test_5351_rf, label='actual')
plt.scatter(X_test_5351_rf['DATEPRD'], y_pred_5351_rf, label='predicted')
plt.xlabel('Year')
plt.ylabel('Bore oil volume')
plt.title('Actual vs Random Forest Regression Model prediction for Bore Oil Volume')
plt.legend()
plt.show()

# Plot residual plot
plt.figure(figsize=(14, 8))
plt.scatter(y_pred_5351_rf, y_pred_5351_rf - y_test_5351_rf, c='b', s=40, alpha=0.5)
plt.hlines(y=0, xmin=0, xmax=max(y_test_5351_rf), colors='r', zorder=3)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot for Random Forest Regression Model for Well No. - 5351')
plt.show()

```






```


In [69]: # Select data for well 7078
X_test_7078_rf = X_test[X_test['NPD_WELL_BORE_CODE'] == 7078]
y_test_7078_rf = y_test[X_test['NPD_WELL_BORE_CODE'] == 7078]
y_pred_7078_rf = rf_model.predict(X_test_7078_rf)

# Plot predicted vs actual oil production
plt.figure(figsize=(14, 8))
plt.scatter(X_test_7078_rf['DATEPRD'], y_test_7078_rf, label='actual')
plt.scatter(X_test_7078_rf['DATEPRD'], y_pred_7078_rf, label='predicted')
plt.xlabel('Year')
plt.ylabel('Bore oil volume')
plt.title('Actual vs Random Forest Regression Model prediction for Bore Oil Volume')
plt.legend()
plt.show()

# Plot residual plot
plt.figure(figsize=(14, 8))
plt.scatter(y_pred_7078_rf, y_pred_7078_rf - y_test_7078_rf, c='b', s=40, alpha=0.5)
plt.hlines(y=0, xmin=0, xmax=max(y_test_7078_rf), colors='r', zorder=3)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot for Random Forest Regression Model for Well No. - 7078')
plt.show()

```






```


In [70]: # Select data for well 7405
X_test_5599_rf = X_test[X_test['NPD_WELL_BORE_CODE'] == 5599]
y_test_5599_rf = y_test[X_test['NPD_WELL_BORE_CODE'] == 5599]
y_pred_5599_rf = rf_model.predict(X_test_5599_rf)

# Plot predicted vs actual oil production
plt.figure(figsize=(14, 8))
plt.scatter(X_test_5599_rf['DATEPRD'], y_test_5599_rf, label='actual')
plt.scatter(X_test_5599_rf['DATEPRD'], y_pred_5599_rf, label='predicted')
plt.xlabel('Year')
plt.ylabel('Bore oil volume')
plt.title('Actual vs Random Forest Regression Model prediction for Bore Oil Volume')
plt.legend()
plt.show()

# Plot residual plot
plt.figure(figsize=(14, 8))
plt.scatter(y_pred_5599_rf, y_pred_5599_rf - y_test_5599_rf, c='b', s=40, alpha=0.5)
plt.hlines(y=0, xmin=0, xmax=max(y_test_5599_rf), colors='r', zorder=3)
plt.xlabel('Predicted Values')
plt.ylabel('Residuals')
plt.title('Residual Plot for Random Forest Regression Model for Well No. - 5599')
plt.show()

```





In []:

In []:

In []:

Comparison of R2 square values

```
In [71]: import matplotlib.pyplot as plt

# Define data
models = ['Linear Regression', 'Polynomial Regression', 'XGBoost Regression',
r2_scores = [0.5545, 0.9301, 0.9941, 0.9755]


# Set figure size
fig, ax = plt.subplots(figsize=(10, 8))

# Define colors for each model
colors = ['blue', 'green', 'red', 'orange']

# Create bar chart
ax.bar(models, r2_scores, color=['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728'])

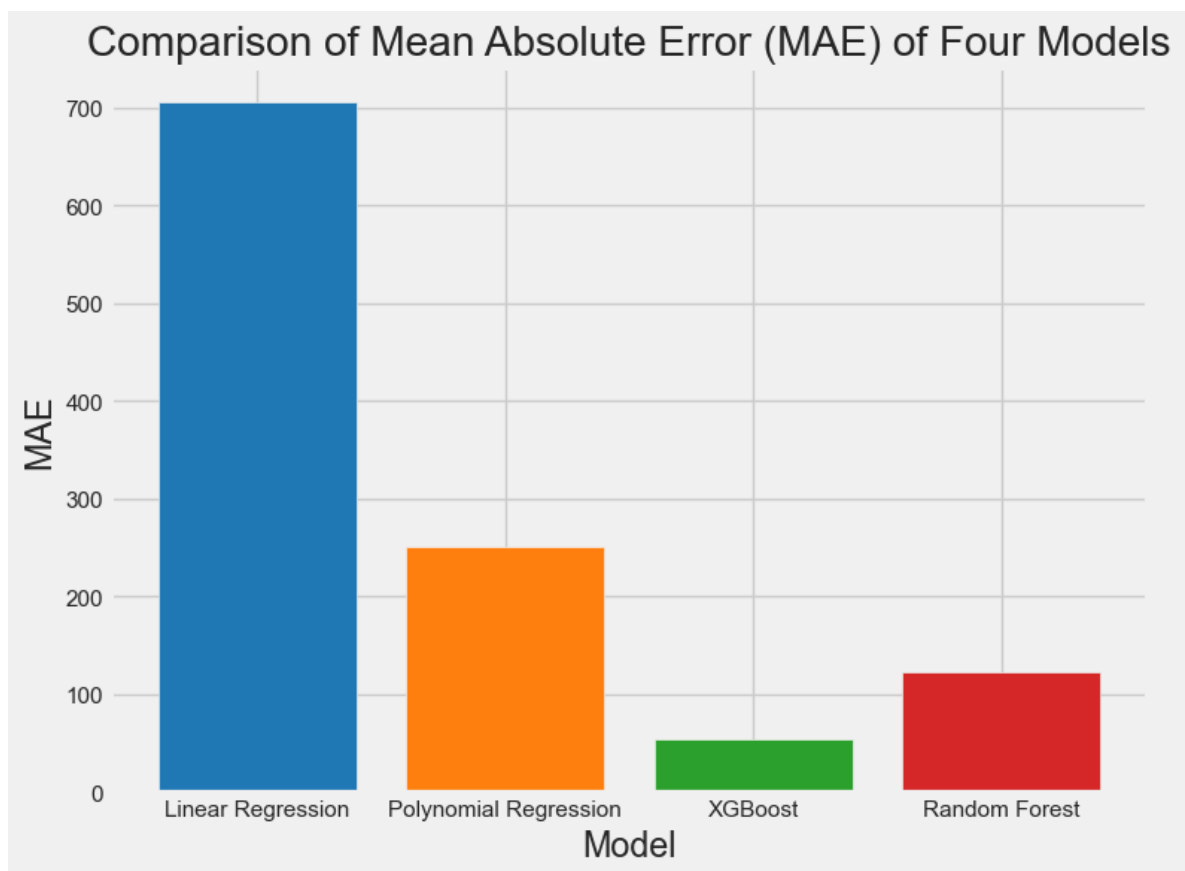
# Add labels and title
ax.set_xlabel('Model')
ax.set_ylabel('R2 Score')
ax.set_title('Comparison of R2 Scores for Oil Production Models')

# Display plot
plt.show()
```



Comparing the Mean Absolute Errors Of the Machine Learning Models

```
In [72]: models = ['Linear Regression', 'Polynomial Regression', 'XGBoost', 'Random Fo  
mae = [705.8759, 250.3786, 53.5111, 122.8333]  
colors = ['blue', 'green', 'red', 'orange']  
fig, ax = plt.subplots(figsize=(8, 6))  
ax.bar(models, mae, color=['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728'])  
ax.set_title('Comparison of Mean Absolute Error (MAE) of Four Models')  
ax.set_xlabel('Model')  
ax.set_ylabel('MAE')  
plt.show()
```




Comparing the Mean Squared Errors of the Machine Learning Models

```
In [73]: import matplotlib.pyplot as plt
import numpy as np

models = ['Linear Regression', 'Polynomial Regression', 'XGBoost', 'Random Fo
mse = [834343.8099, 130862.5409, 11113.8842, 45689.4333]
colors = ['blue', 'green', 'red', 'orange']

fig, ax = plt.subplots(figsize=(8, 6))
ax.bar(models, mse, color=['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728'])
ax.set_title('Comparison of Mean Squared Error (MSE) of Four Models')
ax.set_xlabel('Model')
ax.set_ylabel('MSE')
plt.show()
```




```
In [74]: import matplotlib.pyplot as plt


# define the data
models = ['Linear Regression', 'Polynomial Regression', 'XGBoost Regression',
rmse = [913.4242, 361.7493, 105.4224, 213.7509]

# set the figure size
plt.figure(figsize=(10, 8))

# create the bar chart
plt.bar(models, rmse, color=['#1f77b4', '#ff7f0e', '#2ca02c', '#d62728'])

# add labels and title
plt.xlabel('Model')
plt.ylabel('Root Mean Squared Error')
plt.title('Comparison of RMSE for Oil Volume Production')

# display the plot
plt.show()
```



```

In [75]: import plotly.graph_objects as go
import pandas as pd

# Create a dataframe with the performance metrics of the four models
data = {'Model': ['Linear Regression', 'Polynomial Regression', 'XGBoost Regr',
                 'R2 Score': [0.55, 0.93, 0.99, 0.98],
                 'MAE': [705.88, 250.38, 53.51, 122.83],
                 'MSE': [834343.81, 130862.54, 11113.88, 45689.43],
                 'RMSE': [913.42, 361.75, 105.42, 213.75]}
df = pd.DataFrame(data)

# Normalize the metrics so they can be plotted on the same scale
df_norm = df.drop('Model', axis=1).apply(lambda x: (x - x.min()) / (x.max() -

# Create the radar chart
fig = go.Figure()


# Add a trace for each model
for model in df['Model']:
    fig.add_trace(go.Scatterpolar(
        r=df_norm.loc[df['Model'] == model].values.flatten().tolist(),
        theta=df_norm.columns.tolist(),
        fill='toself',
        name=model
    ))

# Set the title and layout
fig.update_layout(
    polar=dict(
        radialaxis=dict(
            visible=True,
            range=[0, 1]
        )
    ),
    showlegend=True,
    title='Performance Metrics of Four Models',
)

# Show the chart
fig.show()

```

Performance Metrics of Four Models



In []:

In []: